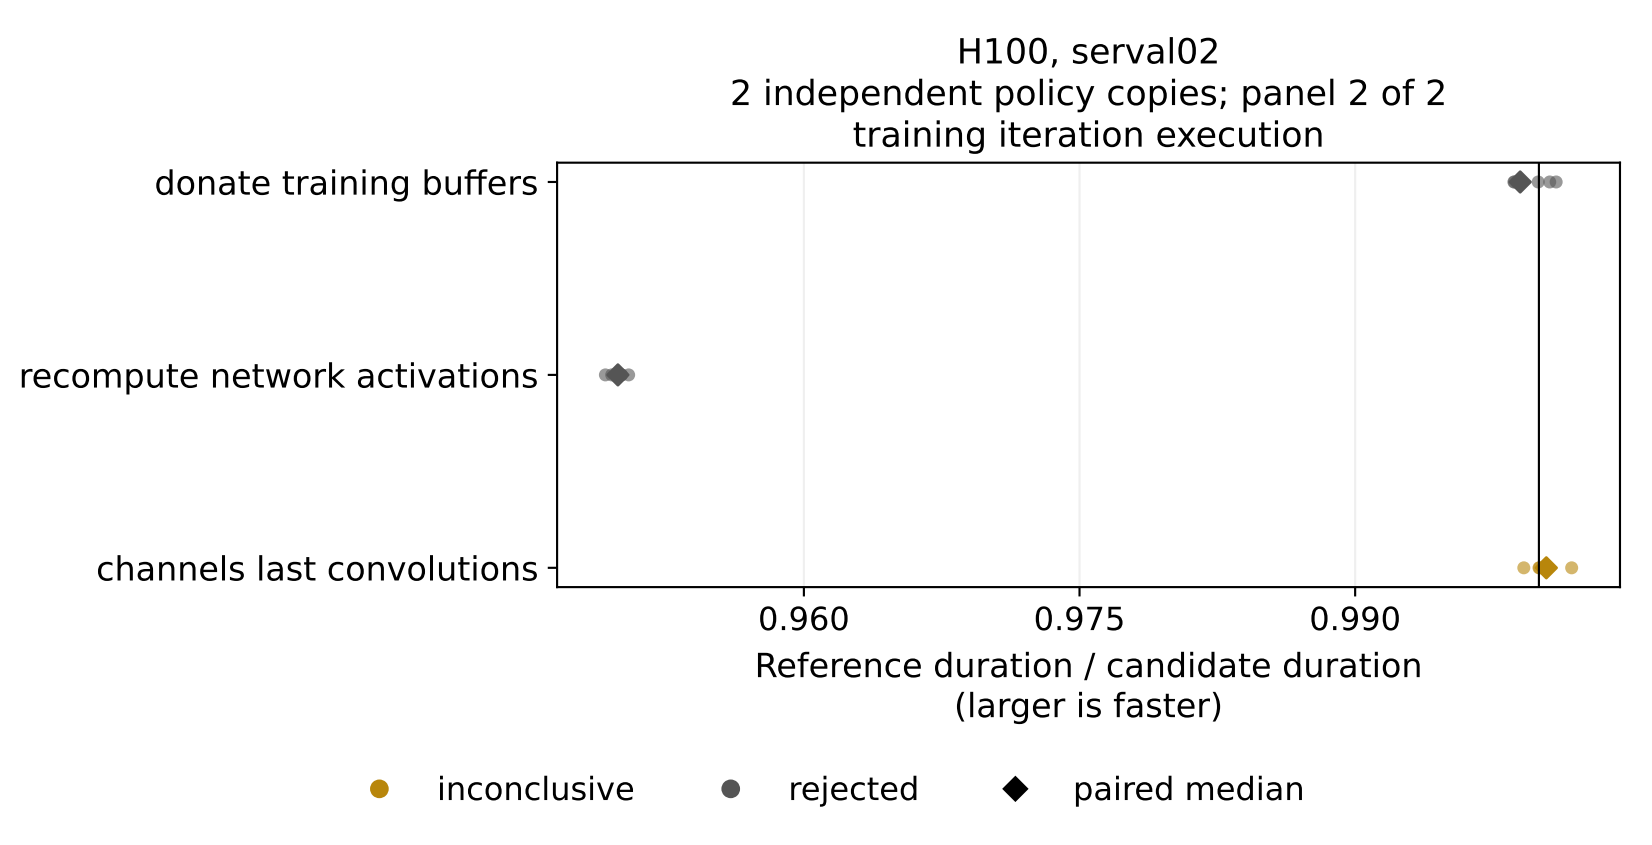
H100, serval02
2 independent policy copies; panel 2 of 2
training iteration execution
donate training buffers
recompute network activations
channels last convolutions
0.960
0.975
0.990
Reference duration / candidate duration
(larger is faster)
inconclusive
rejected
paired median