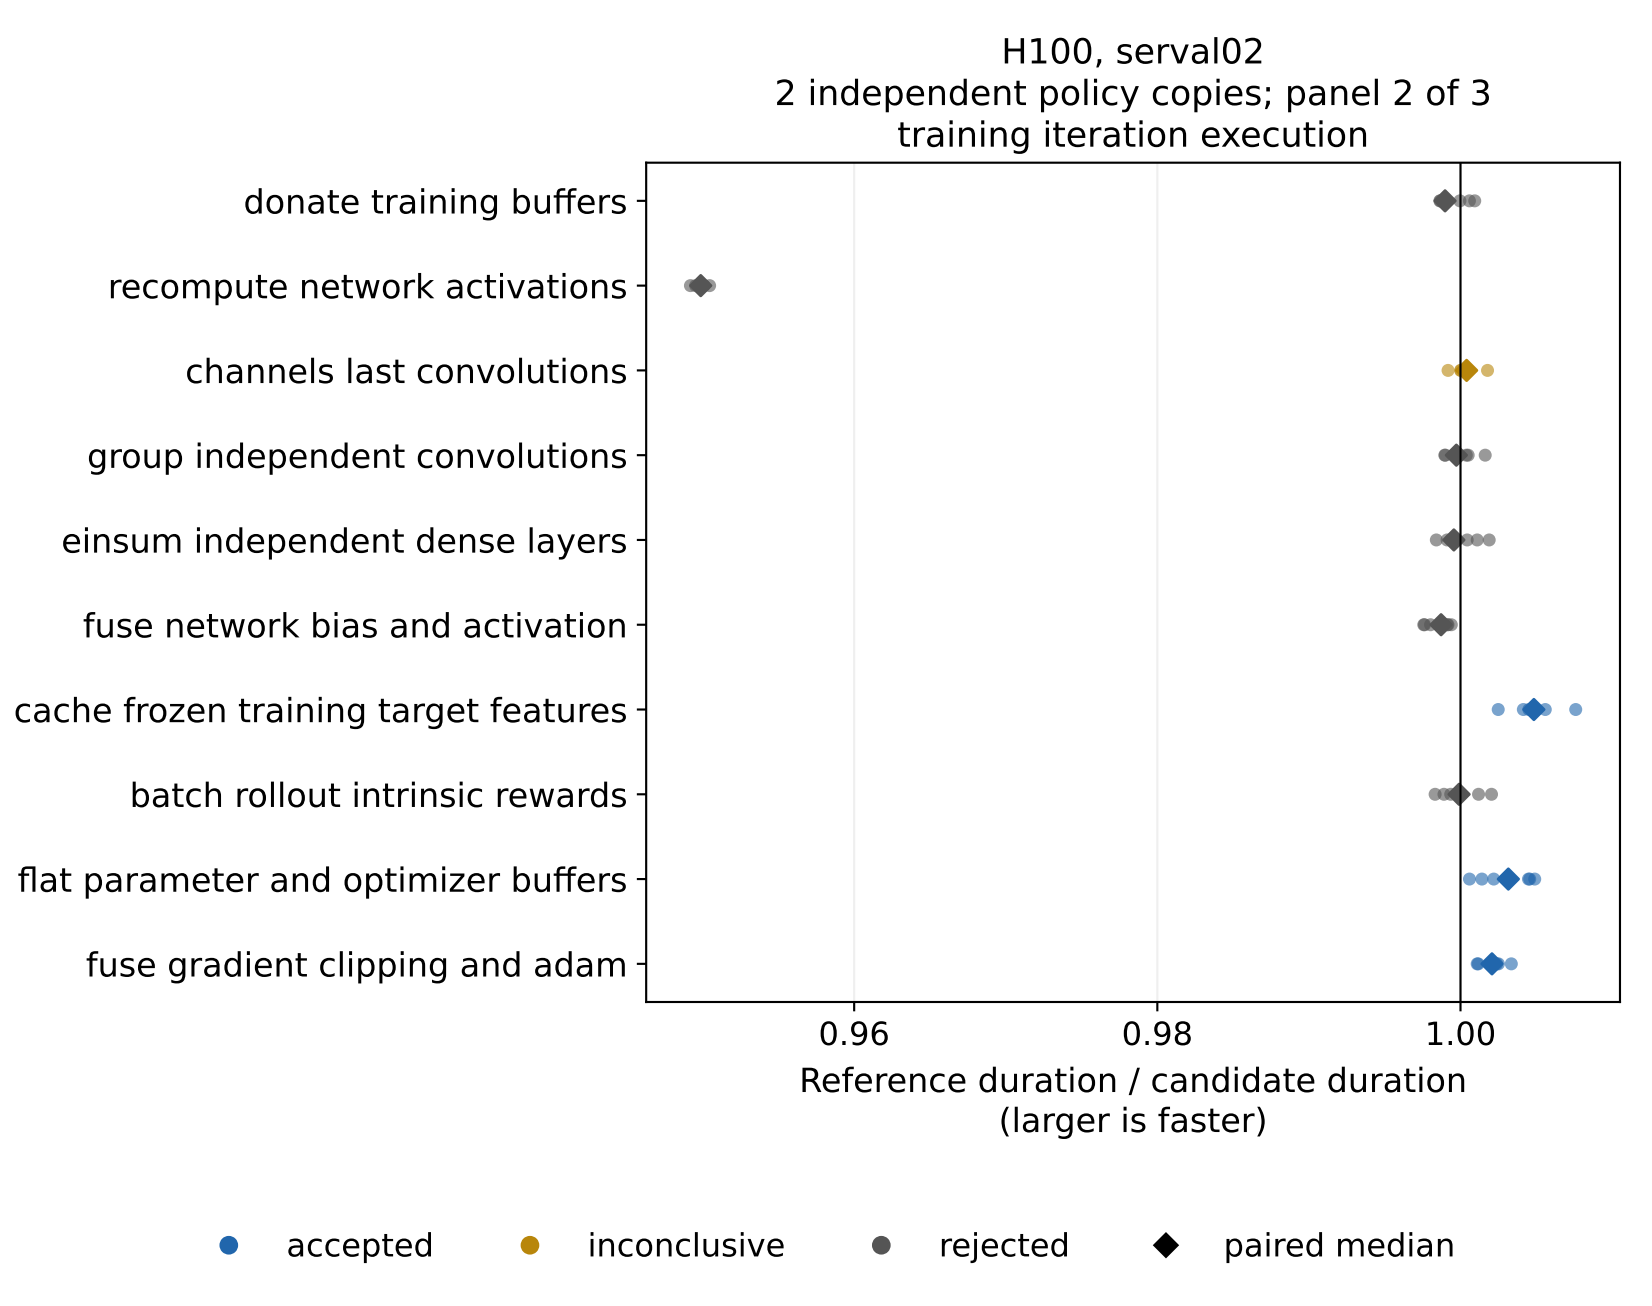
H100, serval02
2 independent policy copies; panel 2 of 3
training iteration execution
donate training buffers
recompute network activations
channels last convolutions
group independent convolutions
einsum independent dense layers
fuse network bias and activation
cache frozen training target features
batch rollout intrinsic rewards
flat parameter and optimizer buffers
fuse gradient clipping and adam
0.96
0.98
1.00
Reference duration / candidate duration
(larger is faster)
accepted
inconclusive
rejected
paired median