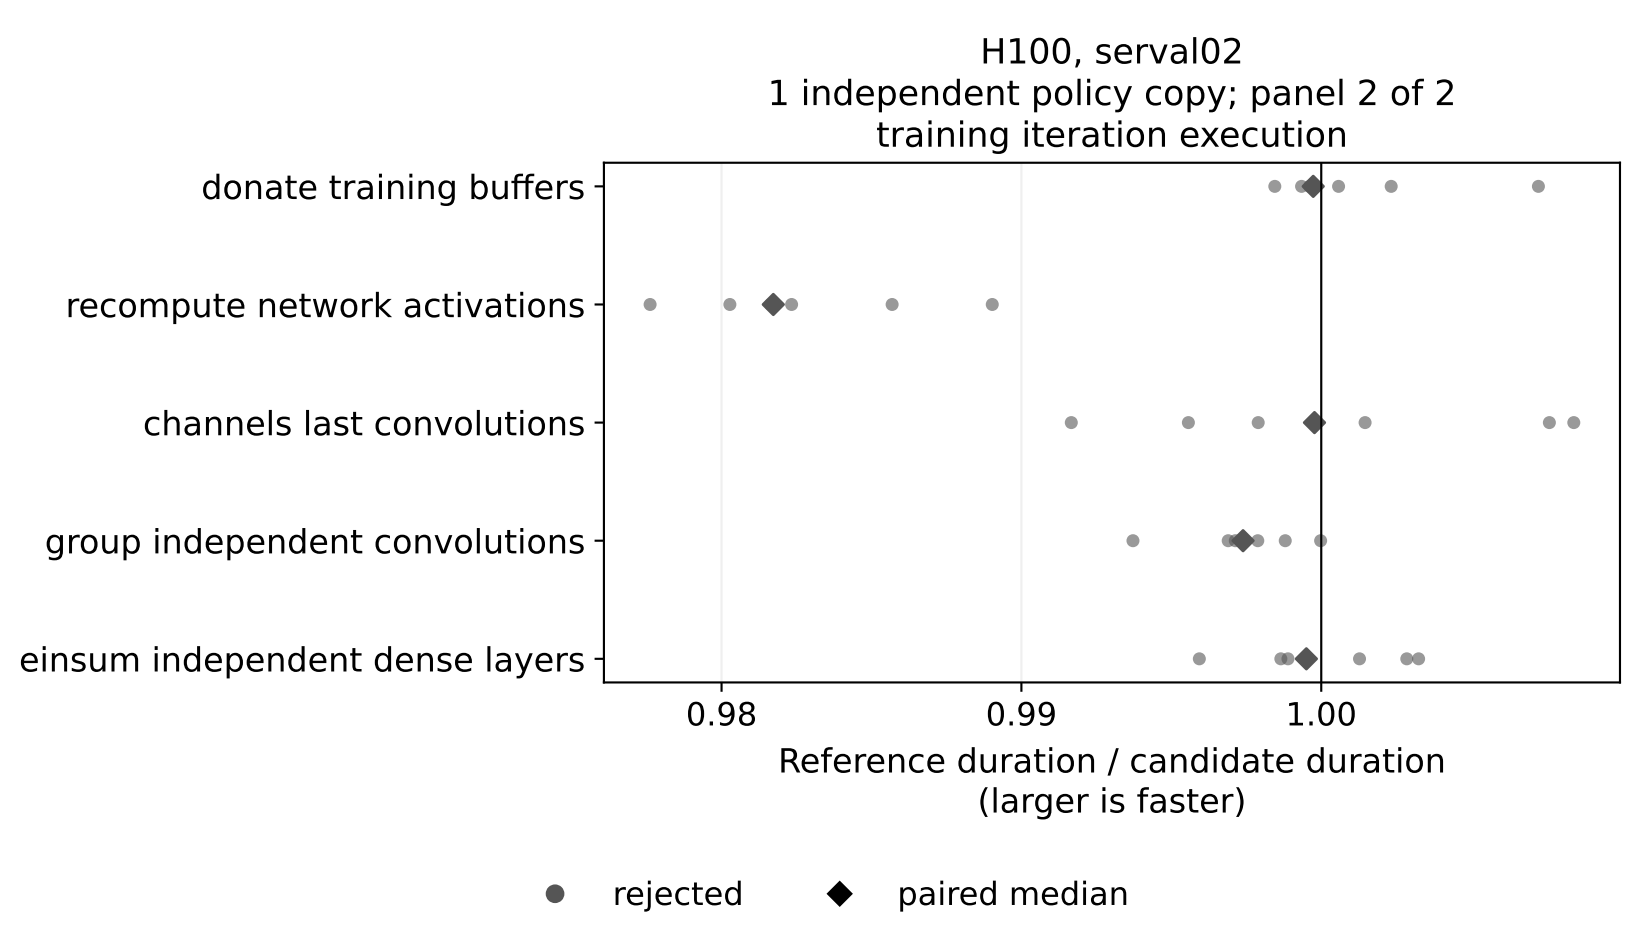
H100, serval02
1 independent policy copy; panel 2 of 2
training iteration execution
donate training buffers
recompute network activations
channels last convolutions
group independent convolutions
einsum independent dense layers
0.98
0.99
1.00
Reference duration / candidate duration
(larger is faster)
rejected
paired median