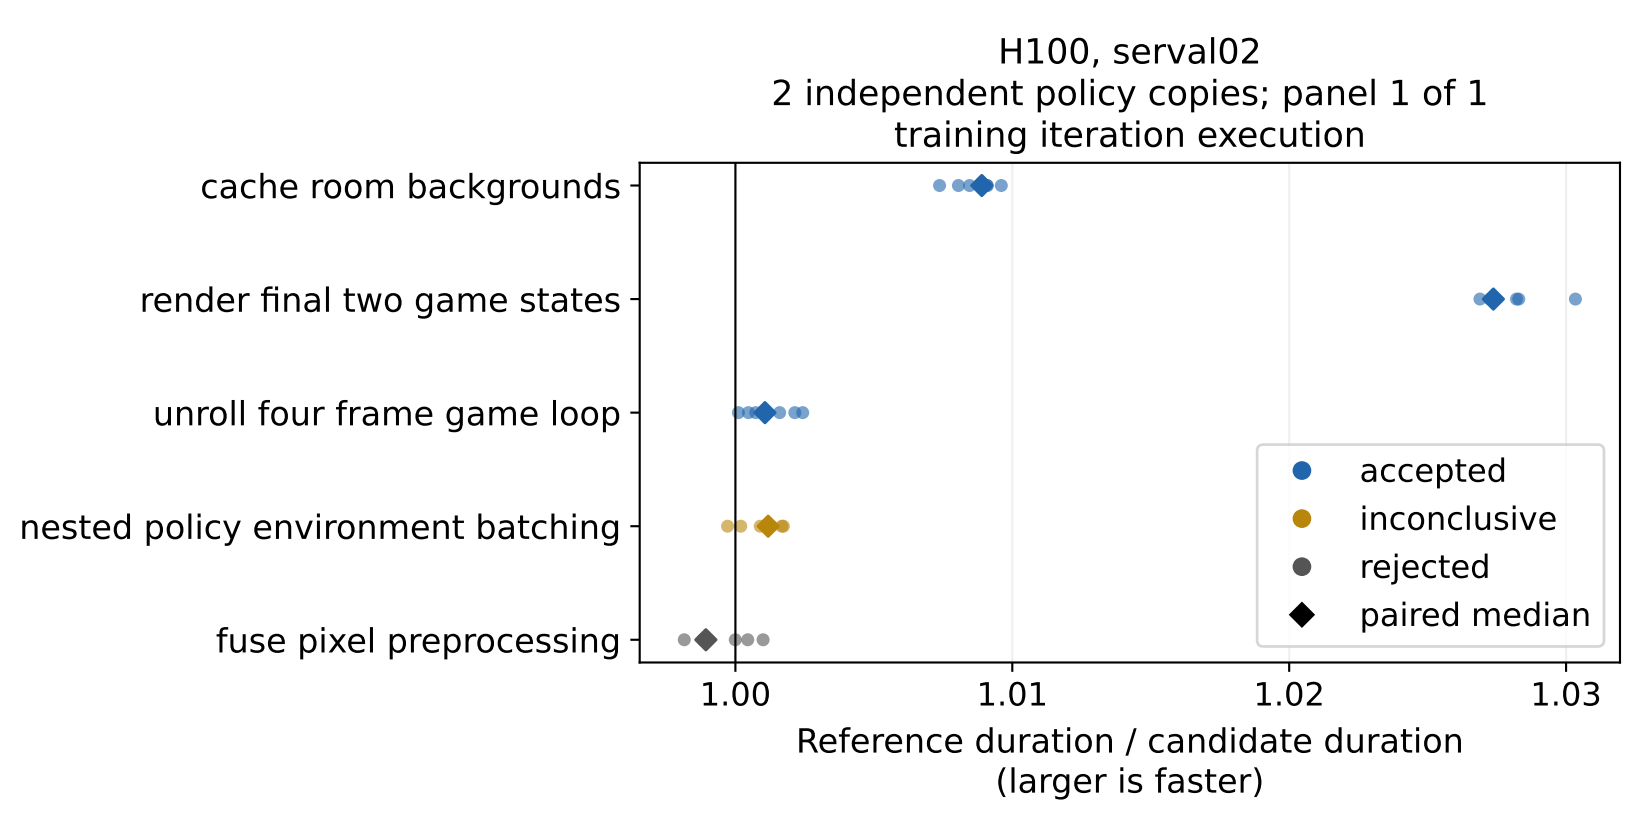
H100, serval02
2 independent policy copies; panel 1 of 1
training iteration execution
cache room backgrounds
render final two game states
unroll four frame game loop
nested policy environment batching
fuse pixel preprocessing
accepted
inconclusive
rejected
paired median
1.00
1.01
1.02
1.03
Reference duration / candidate duration
(larger is faster)