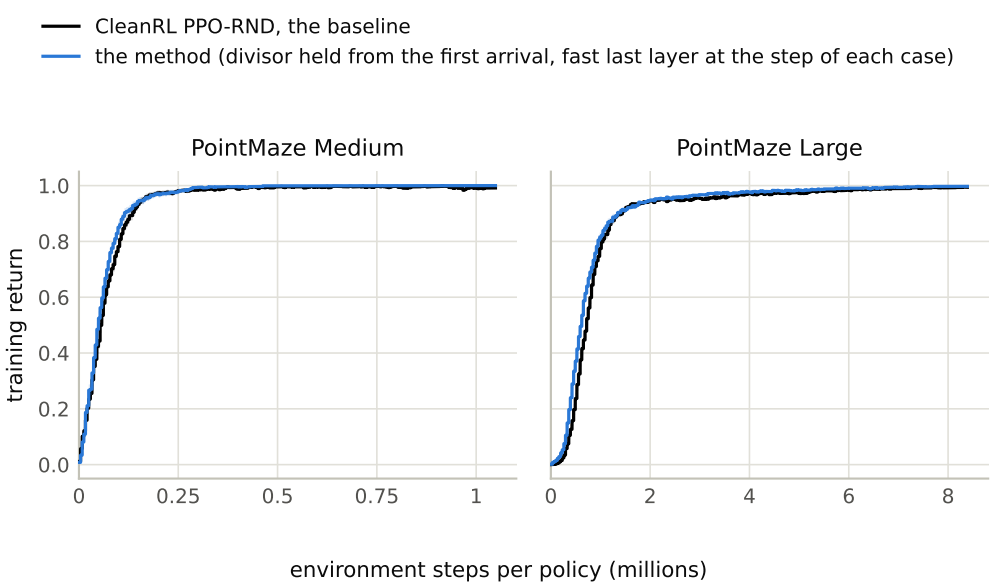
CleanRL PPO-RND, the baseline
the method (divisor held from the first arrival, fast last layer at the step of each case)
PointMaze Medium
PointMaze Large
training return
1.0
0.8
0.6
0.4
0.2
0.0
0
0.25
0.5
0.75
1
0
2
4
6
8
environment steps per policy (millions)