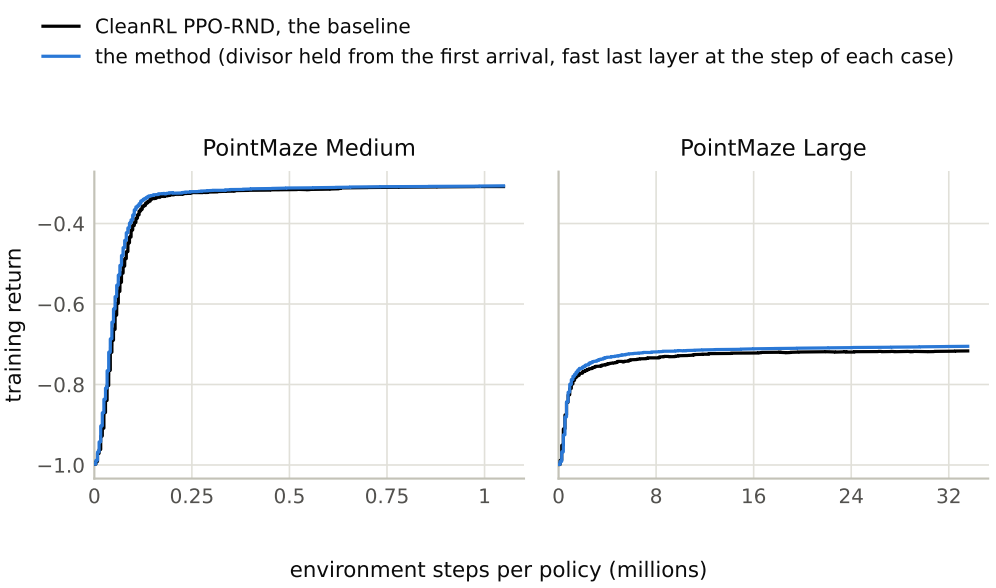
CleanRL PPO-RND, the baseline
the method (divisor held from the first arrival, fast last layer at the step of each case)
PointMaze Medium
PointMaze Large
training return
−0.4
−0.6
−0.8
−1.0
0
0.25
0.5
0.75
1
0
8
16
24
32
environment steps per policy (millions)