Development, straight grid, 32 cells, neural policy, 16 seeds per point (dotted: no bonus at each learning rate and entropy)

uniform field, reward normalization off

uniform field, reward normalization on

alternating high/low field (64:1), reward normalization off

alternating high/low field (64:1), reward normalization on

frozen network field, state normalization off, reward normalization off

frozen network field, state normalization off, reward normalization on

frozen network field, state normalization on, reward normalization off

frozen network field, state normalization on, reward normalization on

online RND, state normalization off, reward normalization off

online RND, state normalization off, reward normalization on

online RND, state normalization on, reward normalization off

online RND, state normalization on, reward normalization on

mean capped first-goal steps

10^7

10^6

10^5

10^4

learning rate 0.0001, entropy 0.001

learning rate 0.0001, entropy 0.01

learning rate 0.001, entropy 0.001

learning rate 0.001, entropy 0.01

2^{-6} 2^{-4} 2^{-2} 2^0 2^2 2^4 2^6

intrinsic weight